Task 2 reflection: AI sycophancy and AI psychosis

An AI assistant can be useful because it responds quickly and speaks with apparent confidence. Those same qualities make agreement especially tempting. If a user proposes an interpretation, a sycophantic assistant may repeat and strengthen it instead of checking the evidence. The user then sees a fluent answer that sounds like independent confirmation, even when the model has merely followed the framing of the question.

I see this risk in investment research. I could start with a belief that a company is well managed, then ask an agent to find reasons to support it. A compliant model could emphasize a favorable shareholder letter, overlook declining margins, and turn an employee's mixed comments into a positive verdict. Repeating that process would produce a neat report without improving the underlying decision. The financial calculations in this homework therefore belong in SQL and Python, and the model's qualitative labor judgment should be tied to the actual interview. A summary should show the evidence and the limits of that evidence, including cases where a tool failed.

The phrase “AI psychosis” raises a more serious concern: some people may come to treat a conversational system as a source of special certainty or personal revelation. I would not assume that an assistant by itself causes a clinical condition. A person's vulnerability, circumstances, and access to support matter. Still, a system that continually validates implausible beliefs could reinforce them. The response style matters: an assistant can acknowledge distress without endorsing an unsupported claim, ask for independent evidence, and encourage real-world support when the situation calls for it.

My practical lesson from building an agent is to separate evidence gathering from persuasion. The agent should use permitted tools, preserve unfavorable facts, compute quantities outside the language model, and make its rule visible. Human review remains necessary when the evidence is ambiguous or the decision has serious consequences. An AI answer should help me inspect my reasoning, not make disagreement disappear.